

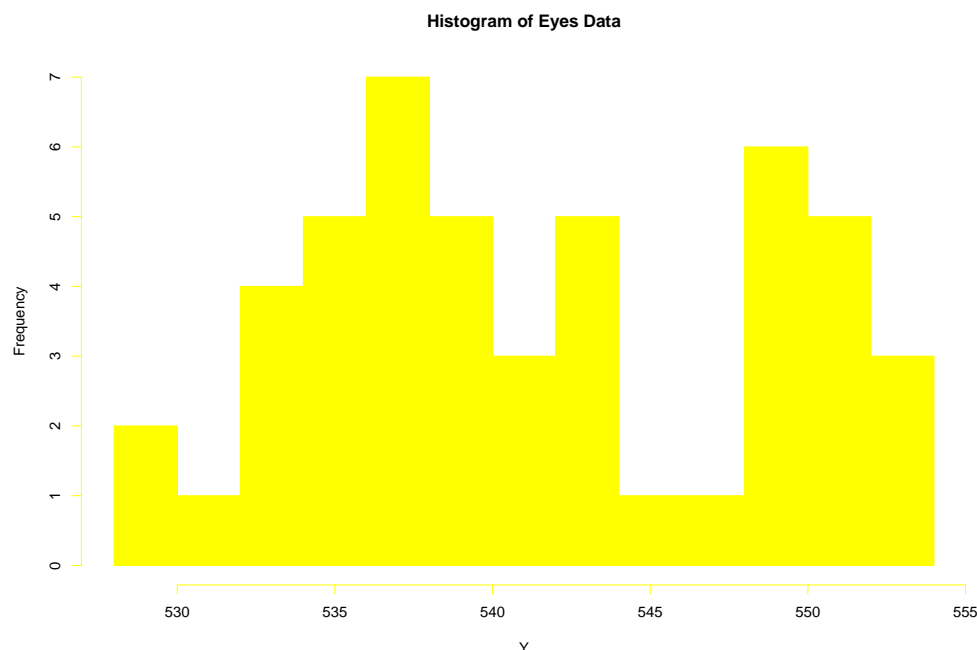
# Mixture Models and Gibbs Sampling

*February 22, 2010*

Readings: Hoff Chapter 6

# Eyes Exmple

Bowmaker et al (1985) analyze data on the peak sensitivity wavelengths for individual microspectrophotometric records on a small set of monkey's eyes. WinBUGs Examples Volume II gives the data for one monkey.



# Mixture Model

Model the data using a Mixture of 2 Normals:

$$Y_i \mid \mu_1, \mu_2, \sigma_1^2, \sigma_2^2, \pi_1, \pi_2 \stackrel{ind}{\sim} \pi_1 \mathbf{N}(\mu_1, \sigma_1^2) + \pi_2 \mathbf{N}(\mu_2, \sigma_2^2)$$

Which is equivalent to

$$Y_i \mid T_i, \mu_1, \mu_2, \sigma_1^2, \sigma_2^2 \stackrel{ind}{\sim} \mathbf{N}(\mu_{T_i}, \sigma_{T_i}^2)$$

$$T_i \stackrel{iid}{\sim} \text{Cat}(T, \pi)$$

where  $T_i$  is a latent variable indicating which group observation  $i$  belongs to i.e.  $T_i \in \{1, 2\}$  and  $\mathbf{P}(T_i = j) = \pi_j$ , and  $\sum_j \pi_j = 1$

# Prior Distributions

Based on WinBUGS example, adopt noninformative prior distributions

$$\mu_j \stackrel{iid}{\sim} \text{N}(0, 1.0 \times 10^6)$$

$$1/\sigma_j^2 \stackrel{iid}{\sim} \text{G}(0.001, 0.001)$$

$$(\pi_1, \pi_2) \sim \text{Dirichlet}(1, 1) \Leftrightarrow \pi_1 \sim \text{Beta}(1, 1)$$

Proper prior distributions are necessary for Mixture Models; if prior on  $\mu$  or  $\sigma^2$  is improper, then the posterior will also be improper if all observations are in one group! False sense of security with vague but proper priors...

# Single Component Gibbs Sampler

Find full conditional distributions for

- $\mu_1 \mid \mu_2, \sigma_1^2, \sigma_2^2, \pi_1, \pi_2, T_1, \dots, T_N, Y$  (normal)
- $\mu_2 \mid \mu_1, \sigma_1^2, \sigma_2^2, \pi_1, \pi_2, T_1, \dots, T_N, Y$  (normal)
- $\sigma_1^2 \mid \mu_1, \mu_2, \sigma_2^2, \pi_1, \pi_2, T_1, \dots, T_N, Y$  (gamma)
- $\sigma_2^2 \mid \mu_1, \mu_2, \sigma_1^2, \pi_1, \pi_2, T_1, \dots, T_N, Y$  (gamma)
- $T_i \mid \mu_1, \mu_2, \sigma_1^2, \sigma_2^2, \pi_1, \pi_2, T_{(i)}, Y$  (Categorical)
- $(\pi_1, \pi_2) \mid \mu_1, \mu_2, \sigma_1^2, \sigma_2^2, T_1, \dots, T_N, Y$  (Dirichlet)

Easy to find and sample!

# Programs

BUGS: Bayesian inference Using Gibbs Sampling

- WinBUGS is the Windows implementation
  - can be called from R with `R2WinBUGS` package
  - can be run on any intel-based computer using VMware, wine
- OpenBUGS open source version of WinBUGS
- LinBUGS is the Linux implementation of OpenBUGS.
- JAGS: Just Another Gibbs Sampler is an alternative program that uses the same model description as BUGS (Linux, MAC OS X, Windows)

Include more than just Gibbs Sampling

# BUGS

Need to specify

- Model
- Data
- Initial values

May do this through ordinary text files or use the functions in `R2WinBUGS` to specify model, data, and initial values then call `WinBUGS`.

# Model Specification via R2WinBUGS

```
mixmodel=function() {  
  for( i in 1 : N ) {  
    y[i] ~ dnorm(mu[i], tau)  
    mu[i] <- lambda[T[i]]  
    T[i] ~ dcat(pi[]) }  
  pi[1:2] ~ ddirch(alpha[])  
  theta ~ dnorm(0.0, 1.0E-6)%_I(0.0, )  
  lambda[1] ~ dnorm(0.0, 1.0E-6)  
  lambda[2] <- lambda[1] + theta  
  tau ~ dgamma(0.001,0.001)  
  sigma <- 1 / sqrt(tau)  
}
```



# Notes on Models

- Distributions of stochastic “nodes” are specified using  $\sim$
- Assignment of deterministic “nodes” uses `<-` (NOT `=`)
- Cannot put expressions as arguments in distributions
- Normal distributions are parameterized using precisions, so `dnorm(0, 1.0E-6)` is a  $N(0, 1.0 \times 10^6)$
- uses `for` loop structure as in R

# Alternative Parameterization

- With vague prior distributions, the Gibbs sampler may get stuck with all observations assigned to one component (hard to escape)
- Label switching Problem
- Robert suggested parameterizing means

$$\lambda_1 \sim \mathcal{N}(0, 1.0 \times 10^6)$$

$$\theta \sim \mathcal{N}_+(0, 1.0 \times 10^6) \quad \theta > 0$$

$$\lambda_2 = \lambda_1 + \theta$$

Constrains Group 2 mean to be larger than Group 1.

# Function to Return Initial Values as a List

```
inits = function() {  
  lambda1 = mean(eyesdata$y[1:30]) + rnorm(1, 0, .  
  theta = mean(eyesdata$y[31:48]) - lambda1  
  sigma2 = var(eyesdata$y[1:30])  
  return(list(lambda = c(lambda1, NA),  
              theta = theta,  
              tau = 1/sigma2,  
              pi = c(30, 48-30)/48))  
}
```

- $\lambda_2$  is not random, so no initial value is specified (it is determined by  $\lambda_1$  and  $\theta$ )
- If no initial value is given, BUGS will generate values given the other values, model and priors

# Data

A list or rectangular data structure for all data and summaries of data used in the model

```
eyesdata= list(  
  y = c(529.0, 530.0, 532.0, 533.1, 533.4, ...  
        535.3, 535.4, 535.9, 536.1, 536.3, 536.4, .  
        538.3, 538.5, 538.6, 539.4, 539.6, 540.4, .  
        543.5, 543.8, 543.9, 545.3, 546.2, 548.8, .  
        549.9, 550.6, 551.2, 551.4, ... 552.9, 553.  
  N = 48,  
  alpha = c(1, 1),  
  T = c(1, NA, NA, NA, NA, NA, NA, NA, NA, ...  
        NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, ...  
        NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, ...  
        NA, 2) )
```

# Notes

- The variable  $T$  is treated as part of the data, rather than “prior”
- With the data sorted, assign the smallest observation to group 1, and the largest to group 2.
- any fixed hyperparameters can be given here

# Specifying which Parameters to Save

The parameters to be monitored and returned to R are specified with the variable `parameters`

```
parameters = c("lambda", "theta", "sigma",  
               "pi" )
```

- To save a whole vector (for example all lambdas, just give the vector name)
- May save stochastic or deterministic nodes

# Running WinBUGS from R

Write the model out as a text file, then call `bugs ( )`

```
path = getwd( )
model.file = paste(path, "model.txt", sep=" ")
write.model(mixmodel, model.file)

sim = bugs(eyesdata, inits, parameters, model.f
          n.chains=2, n.iter=5000,
          bugs.dir=BUGS.DIR, # for use with MA
          WINE=WINE,         #for use with MAC
          WINEPATH=WINEPATH, #for use with MAC
          debug=T, DIC=F)
```

`debug=T` keeps WinBUGS open – very useful for debugging BUGS!

# Output

```
> sim
```

```
  2 chains, each with 5000 iterations  
(first 2500 discarded), n.thin = 5  
n.sims = 1000 iterations saved
```

	mean	sd	2.5%	50%	97.5%	Rhat	n.e
lambda[1]	536.7	0.9	535.0	536.7	538.6	1	10
lambda[2]	548.9	1.2	546.3	548.9	551.3	1	7
theta	12.1	1.4	9.2	12.3	14.6	1	10
pi[1]	0.6	0.1	0.4	0.6	0.8	1	10
pi[2]	0.4	0.1	0.2	0.4	0.6	1	10
sigma	3.8	0.6	3.0	3.6	5.3	1	10